Linux, day 10

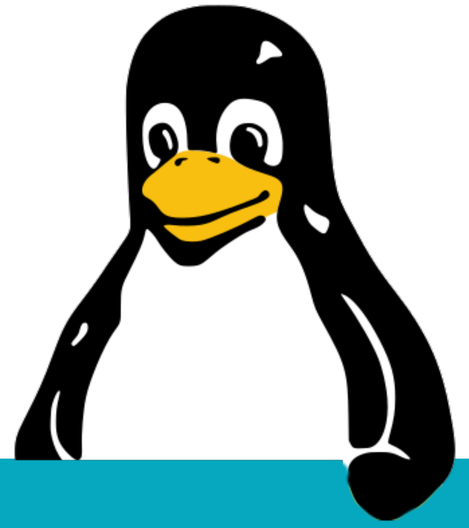

MDadm

Prepping your lab: VBox

- In VirtualBox, change the three dummy disks.
  - Enable hot-plugging.

Prepping your lab: VBox

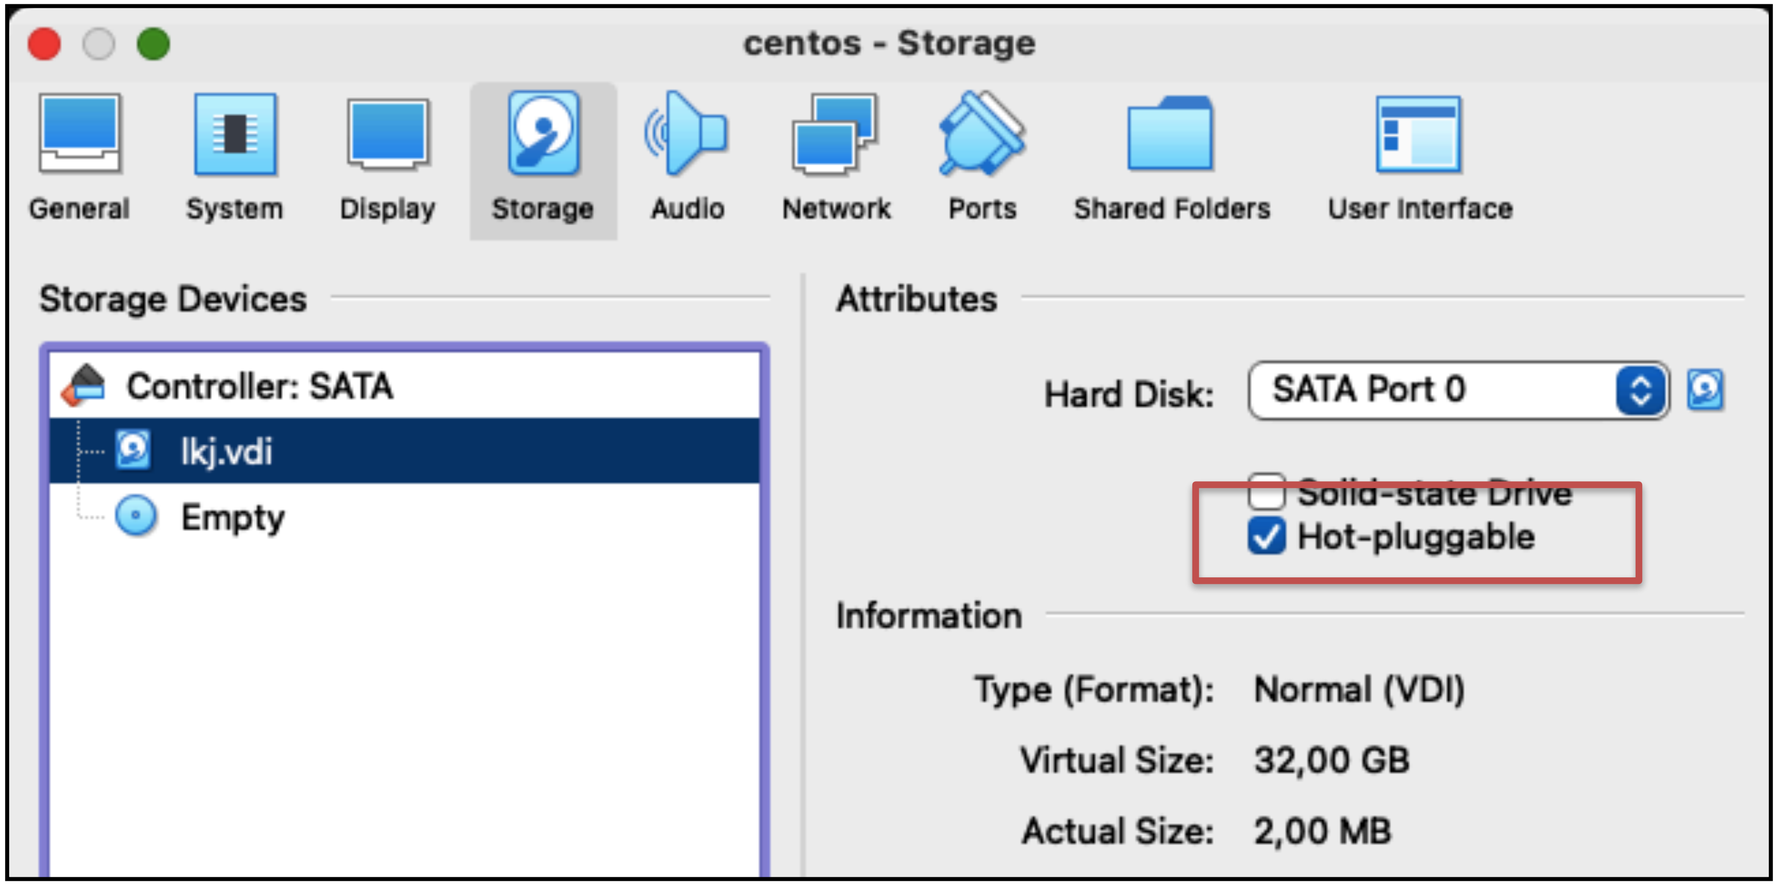

Prepping your lab: mounts

- In the Linux VM verify the disks are unused.
  - *mount | grep "sd?"*
    - Unmount your mounts of *sdb*, *sdc* and *sdd*.
  - *grep -i "sd?" /etc/fstab*
    - Remove any lines for *sdb*, *sdc* and *sdd*.
  - *sudo mkdir /mnt/data*

Prepping disks for RAID

- Some people suggest you partition them,
  - Others say to just use the whole device.
  - *mdadm* says "partition table will be meaningless"
- Apply GPT partitioning.
  - Use partition type "*Linux RAID*" (29) in *fdisk.*
  - Or *"FD00"* in *gdisk.*

Prepping our disks (three times)

```
$ sudo fdisk /dev/sdb
g
n    # Three times enter, y to overwrite
t
29   # Type RAID
w
```

Making a RAID1, with spare

- Remember how to use those backslashes?

```
$ sudo mdadm --create --verbose \
 --level=1 --metadata=1.2 \
 --raid-devices=2 /dev/md/MyRAID1 \
 /dev/sdb /dev/sdc \
 --spare-devices=1 /dev/sdd
```

Then use it

- Format, mount and use.

```
$ sudo mkfs.ext4 /dev/md/MyRAID1

$ sudo mount /dev/md/MyRAID1 /mnt/data

$ sudo touch /mnt/data/testfile
```

Saving your array config

- The location of the config file differs per Linux!
  - *sudo find /etc -name "mdadm.conf"*

```
$ sudo mdadm --detail --scan | \
sudo tee -a /etc/mdadm.conf
                          # Change that path!
```

Checking RAID health

- *mdadm* software RAID is pretty basic!
  - *cat /proc/mdstat*
  - *sudo mdadm --detail /dev/md/MyRAID1*
  - *journalctl | grep -e "kernel: md" -e "mdadm"*

Fun time

- On the guest OS, keep an eye on RAID.

```
$ sudo mdadm --monitor /dev/md/MyRAID1
```

Let's break it!

- Remove the second 100MB virtual disk.

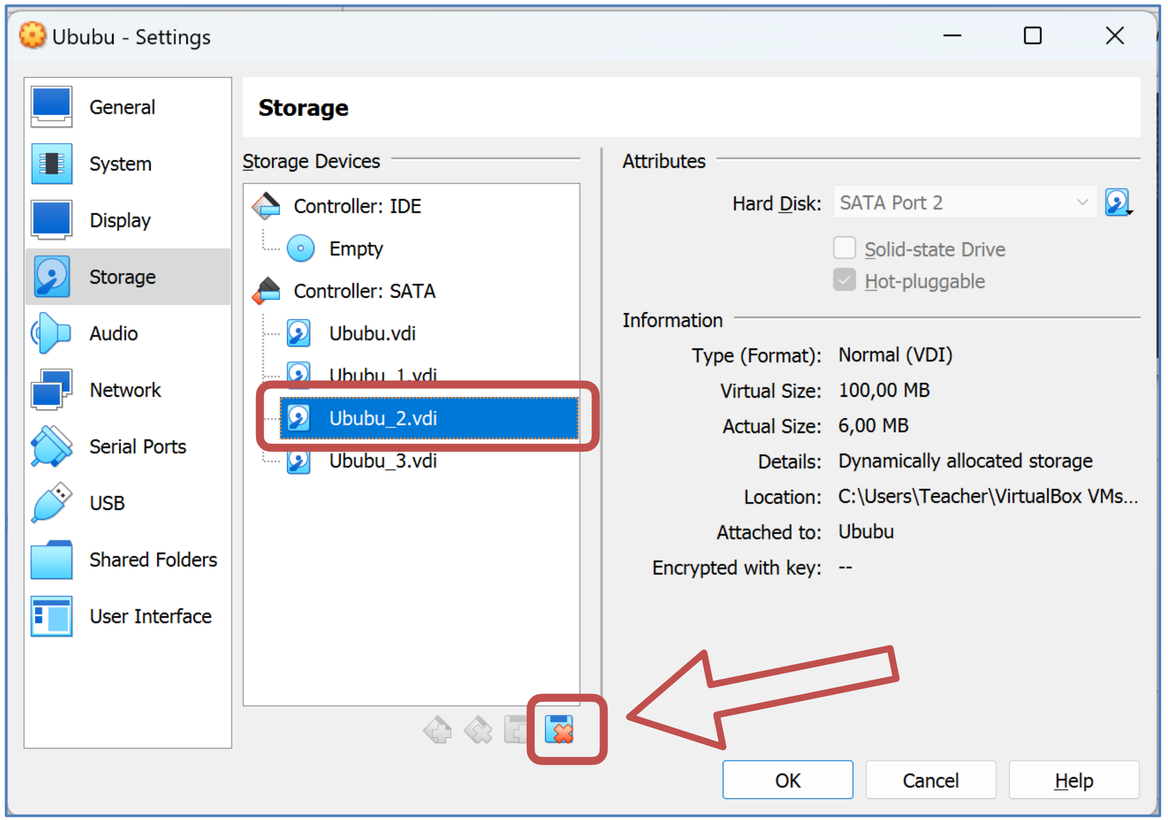

What happened?

- *mdadm* noticed the failure,
  - Grabbed the hot spare and started rebuilding.
  - The failed disk disappeared.

- Check "*/proc/mdstat*", "*mdadm*" and "*journalctl*".

Question!?

- What would be a nice way to test availability?
- If your boss wants you to prove *"mdadm*" is good,
  - How would you run the tests?
  - Can Python or shell scripts help?

Let's rebuild

- Re-attach the virtual disk you removed.

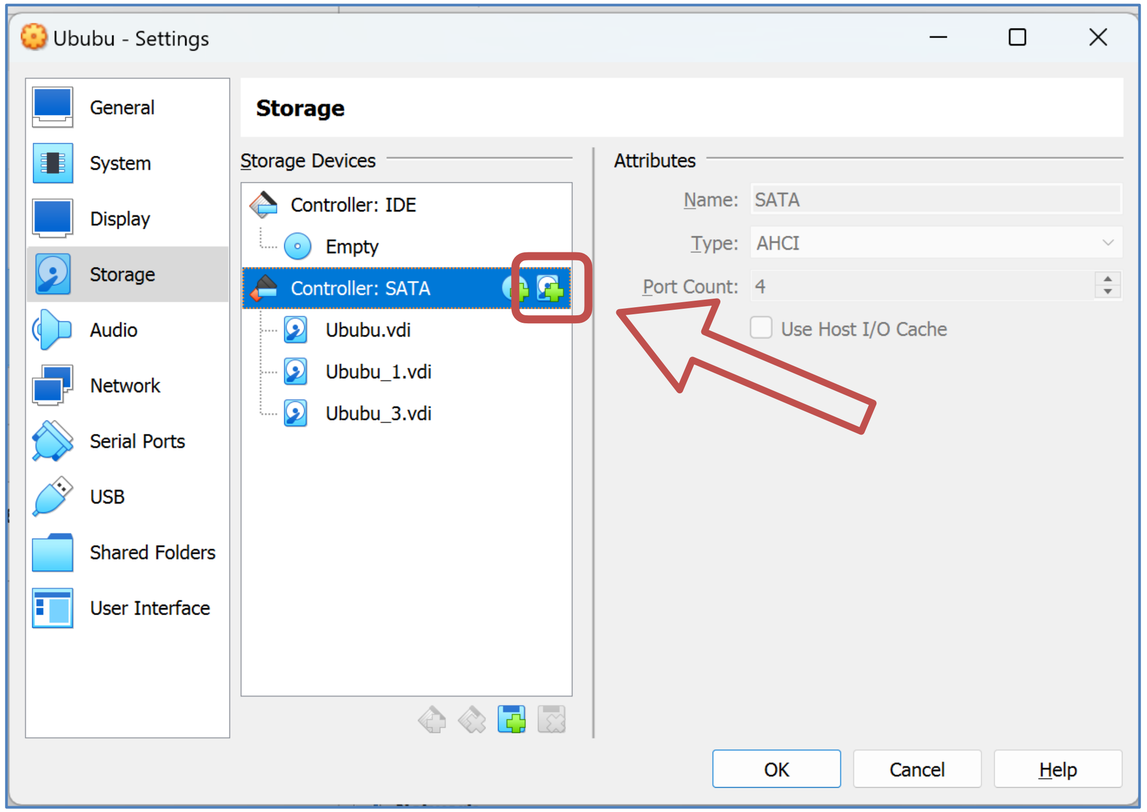

Let's rebuild

- And we will re-add the disk to the RAID set.

```
$ sudo mdadm --manage /dev/md/MyRAID1 \
  --add /dev/sdc

$ sudo cat /proc/mdstat
```

Done? Break it all.

```
$ sudo umount /mnt/data
$ sudo mdadm --stop /dev/md/MyRAID1
$ sudo rm /etc/mdadm.conf  # adjust path!

$ sudo mdadm --zero-superblock /dev/sdb
$ sudo mdadm --zero-superblock /dev/sdc
$ sudo mdadm --zero-superblock /dev/sdd
```

For your homework: RAID5

- Needs three virtual disks!

```
$ sudo mdadm --create --verbose \
 --level=5 --metadata=1.2 \
 --raid-devices=3 /dev/md/MyRAID5 \
 /dev/sdb /dev/sdc /dev/sdd
```

LVM

The problem we need to solve

- Partitioning disks allows some flexibility.
- But growing, shrinking and moving?
  - That's a huge hassle!

- LVM allows much greater flexibility,
  - Through another layer of virtualization.

LVM concepts

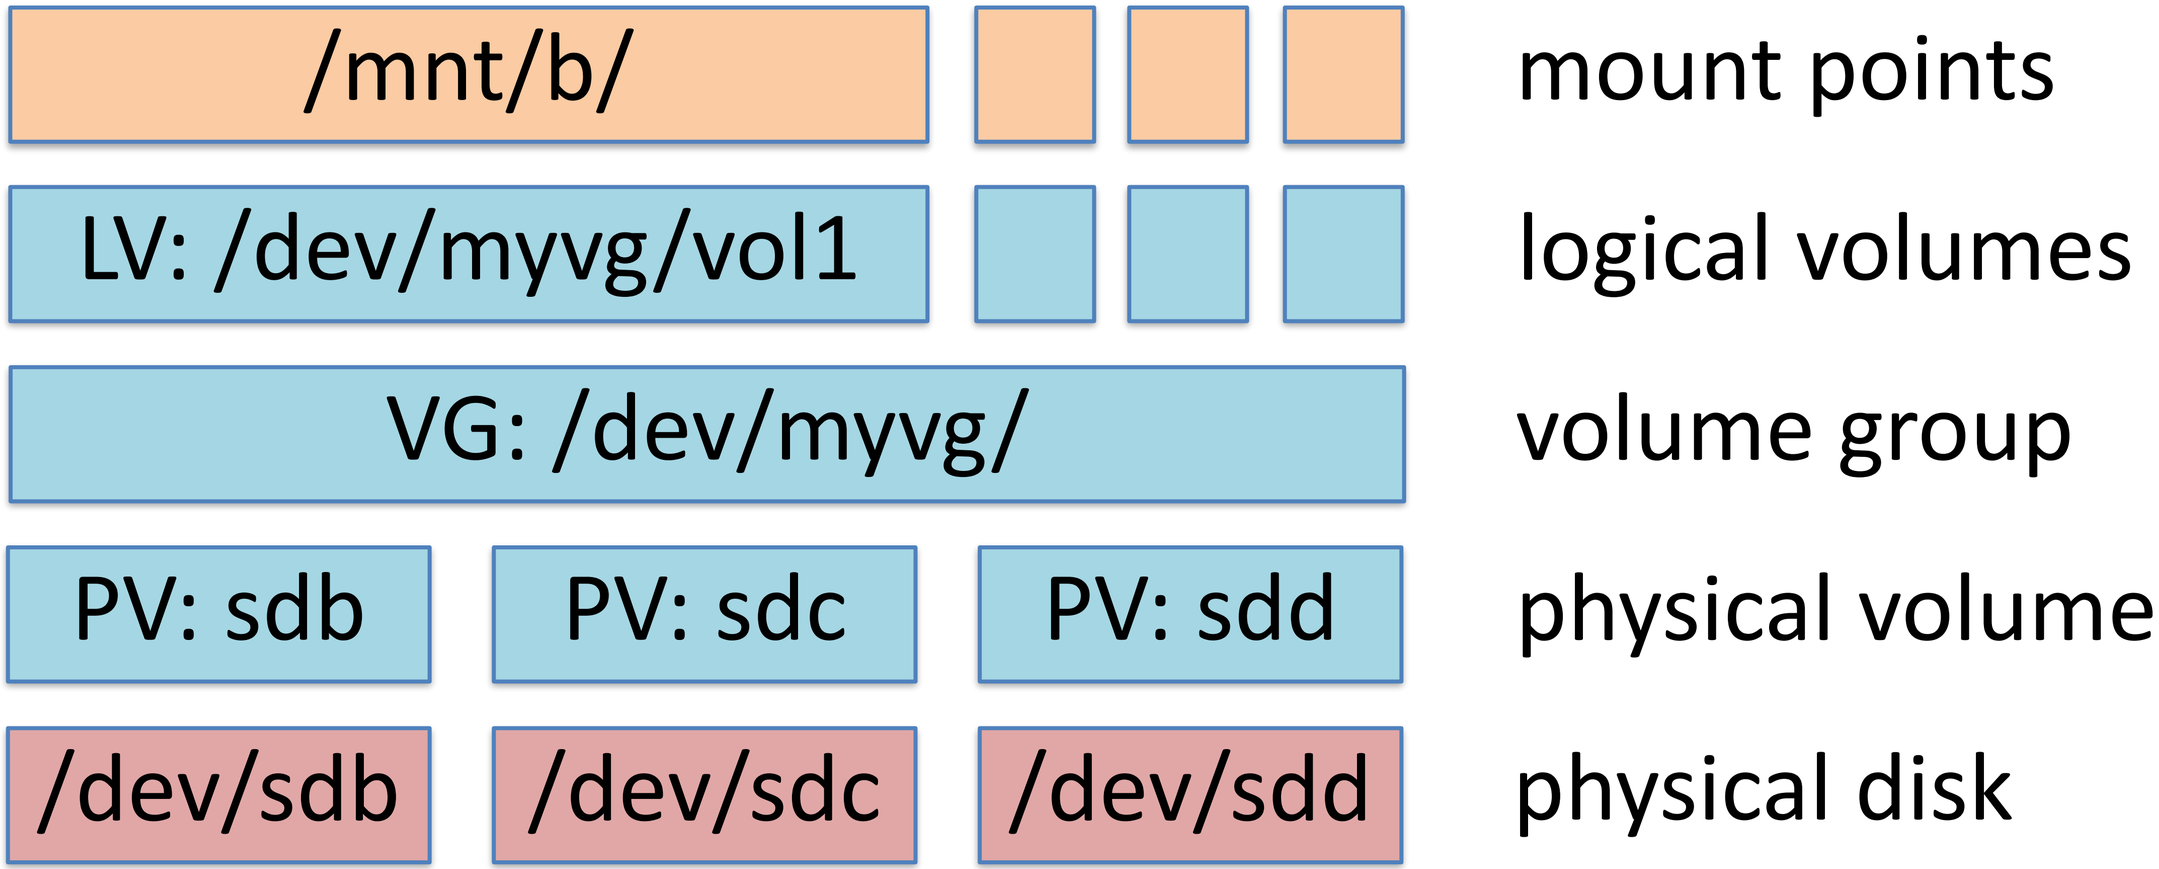

Prepping our disks (three times)

```
$ sudo fdisk /dev/sdb
g
n    # Three times enter, y to overwrite
t
30   # Type LVM
```

Creating phys. volumes

- This adds the device to LVM control

```
$ sudo pvcreate /dev/sdb1
$ sudo pvcreate /dev/sdc1
$ sudo pvcreate /dev/sdd1
```

That may fail...

- "*pvcreate*" will notice if disks are re-used.
  - Wipe and try again!

```
$ sudo wipefs -a /dev/sdb
$ sudo wipefs -a /dev/sdc
$ sudo wipefs -a /dev/sdd
```

Making the volume group

- This collects the disks for usage.

```
$ sudo vgcreate myvg \
/dev/sdb1 /dev/sdc1 /dev/sdd1

$ sudo vgdisplay myvg
```

Our first volume!

- This collects the disks for usage.

```
$ sudo lvcreate -L 50M -n vol1 myvg

$ sudo vgdisplay
$ sudo lvdisplay
```

Then use it

- Format, mount and use.

```
$ sudo mkfs.ext4 /dev/myvg/vol1

$ sudo mount /dev/myvg/vol1 /mnt/data

$ sudo touch /mnt/data/testfile
```

Check things out

- The commands start with pv*, vg*, lv*.

pvdisplay	Show all physical volumes
vgdisplay	Show all volume groups
lvdisplay	Show all logical volumes

Resizing file systems

- For example, on a volume that was 50 MB.

```
$ sudo lvresize -L 70M /dev/myvg/vol1

$ sudo resize2fs /dev/myvg/vol1
```

Tear it all down!

- Breaking what we made is easy.

```
$ sudo umount /mnt/data
$ sudo lvremove /dev/myvg/vol1
$ sudo vgremove /dev/myvg
$ sudo pvremove /dev/sdb1 /dev/sdc1 \
 /dev/sdd1
```

Want more?

- You can have multiple volume groups.
  - Each with multiple volumes.
- You can add more disks.
- You can grow existing groups and volumes.

What will we do today?

- ~~Recap~~
- ~~(More) storage management~~
  - ~~Devices, MDadm, LVM~~
- TLS, cryptography and certificates
- Extra lab: NFS server
- Closing: homework and Q&A

LAB: NFS server

Time for research

- Can you make the following?
  - Setup your Ubuntu VM as NFS server.
  - Make /var/userdata on Ubuntu.
  - Make that directory an NFS file share.
  - Setup Fedora as NFS client, and mount the share.

Closing

Homework

- Reading:
  - Chapter 12
  - Chapter 21
  - Chapter 23
  - Chapter 26

Homework

- Go do:
  - Retry your RAID1 setup, incl. the "failed" device.
  - Use your three disks to make one *mdadm* RAID5.
  - On the RAID5 set, make an LVM volume group
    - And a 100MB logical volume, on */mnt/data*.

Homework

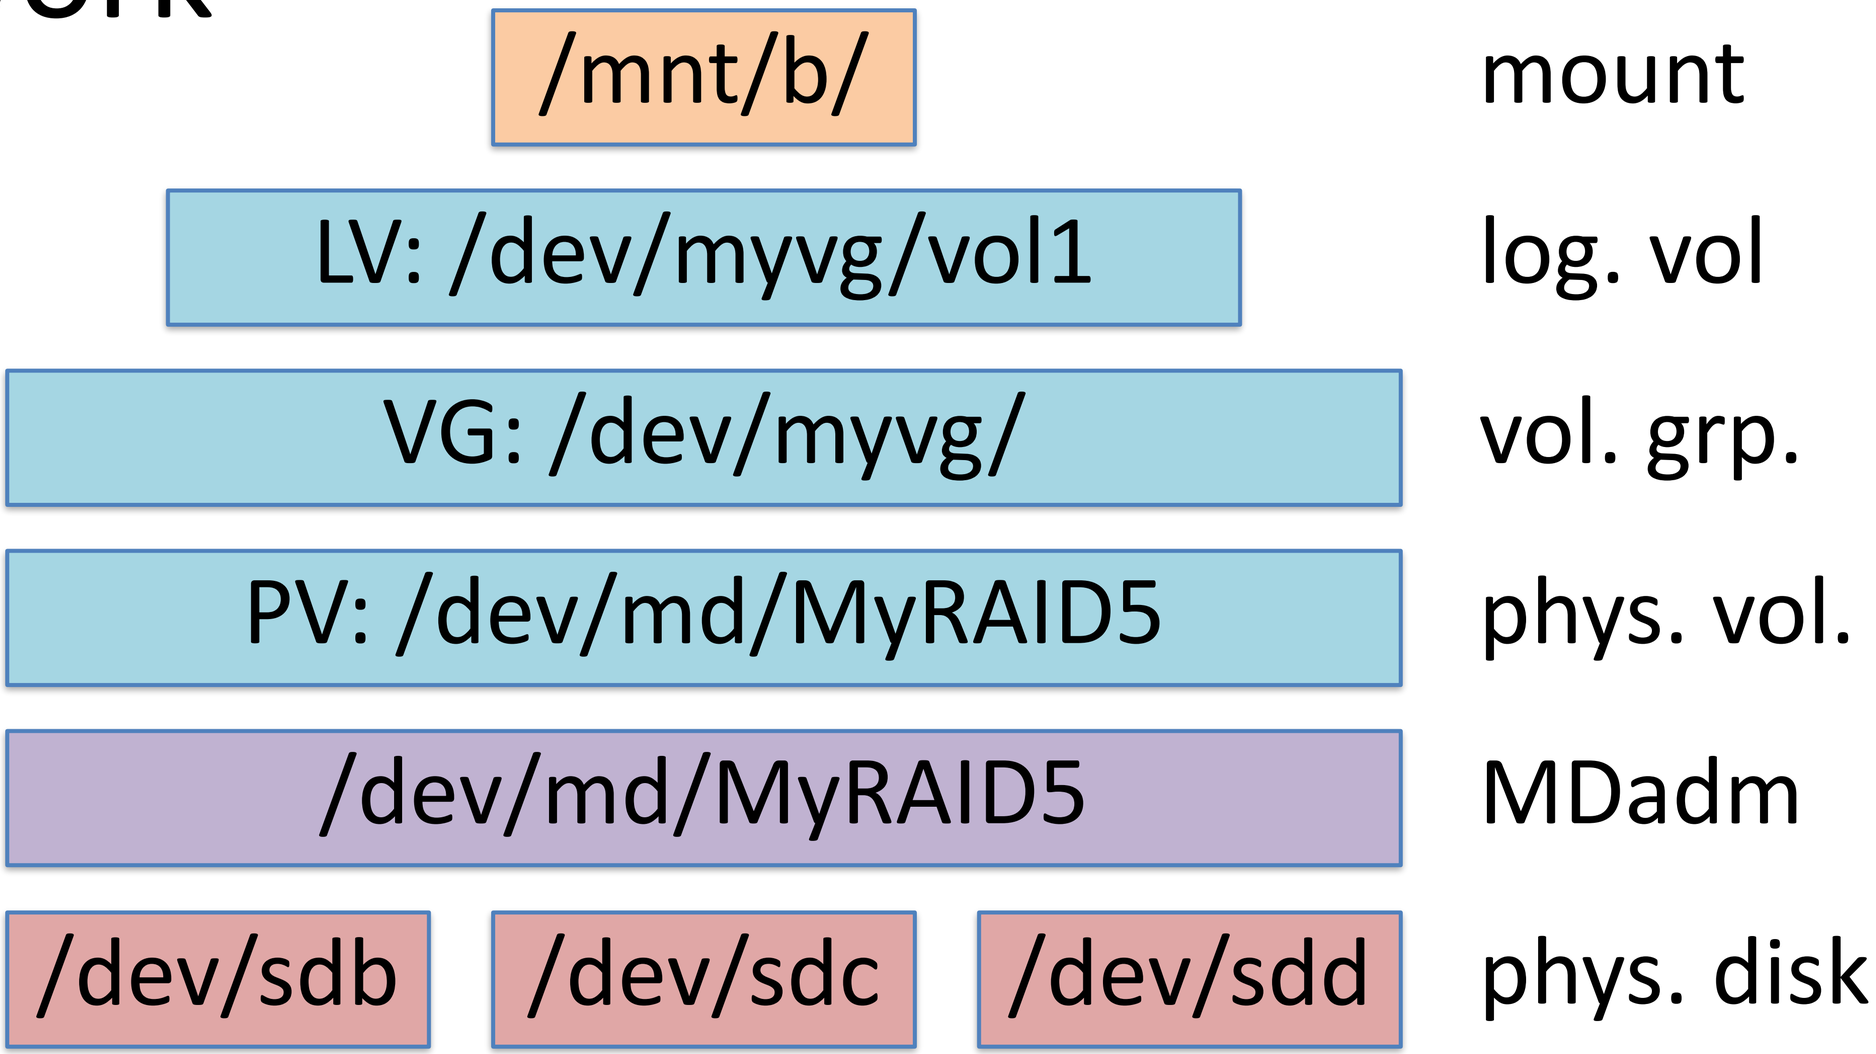
/mnt/b/
mount
LV: /dev/myvg/vol1
log. vol
VG: /dev/myvg/
vol. grp.
PV: /dev/md/MyRAID5
phys. vol.
/dev/md/MyRAID5
MDadm
/dev/sdb
/dev/sdc
/dev/sdd
phys. disk

Advanced homework

- Go do:
  - Take the 100MB logical volume you made.
    - And expand it to 150MB.
    - You need to grow both the logical volume,
    - as well as the file system.